

Introduction

Micro aerial robots, such as small quadrotors, are increasingly tasked with vision-based autonomy like surveillance, search-and-rescue, and navigation. These tasks demand **real-time object detection and multi-object tracking (MOT)** using onboard cameras, but achieving this on lightweight, battery-powered hardware (e.g. Raspberry Pi or NVIDIA Jetson Nano) is challenging. In the past five years (2020–2025), significant research has focused on **efficient computer vision models and systems optimized for deployment on embedded UAV platforms**. Key challenges include the limited computing resources and power on micro drones, the need to maintain high detection accuracy for small, distant objects, and integration of perception outputs with flight control in real time ¹ ². This report reviews recent advances in onboard object detection and MOT for micro aerial robots, emphasizing **lightweight models, hardware acceleration techniques, onboard perception-control integration**, performance benchmarks, and open-source resources.

Lightweight Object Detection Models for Embedded UAVs

Real-time **object detection** on micro UAVs has gravitated toward one-stage CNN models that balance accuracy and speed. The YOLO family (You Only Look Once) has been particularly popular, with “tiny” or nano versions designed for efficiency. For example, *YOLOv5-Nano* and *YOLOv7-Tiny* models (few million parameters) have been demonstrated on embedded GPUs ³. In tests, a YOLOv7-tiny detector on a Raspberry Pi 4 could only achieve ~0.9 FPS, whereas on a Jetson Nano it reached ~5–7 FPS, and up to ~30 FPS on the more powerful Jetson Xavier NX ³. Newer YOLO versions continue this trend: a 2024 study reported that YOLOv8-Nano can run at ~60 FPS on a Jetson Orin Nano (using TensorRT optimization), whereas larger YOLOv8 models (e.g. Medium) run slower (~20 FPS) but with higher accuracy ⁴. These results highlight the trade-off between model size and real-time performance on edge devices ¹.

Researchers have also explored **alternative lightweight detectors**. One 2023 study modified the *CenterNet* one-stage detector for small object detection on drones ⁵ ⁶. The authors tested CNN backbones from ResNet18 up to Res2Net101 and hourglass networks, comparing them against YOLOv1–v7, SSD-MobileNet and others on aerial datasets (VisDrone, UAV datasets). Their optimized CenterNet (with Res2Net101 backbone) achieved the highest mAP, especially on small objects, outperforming YOLOv7 in accuracy on drone-specific data ⁷ ⁸. However, on an embedded Jetson Nano, even the fastest configuration (ResNet18 backbone) only ran ~4.5 FPS ⁹ ¹⁰, underscoring how heavy models must be pared down for real-time use. The authors note an “acceptable FPS for a drone would be at least 25–30” for safe autonomous flight ², a target not met by most unaccelerated models on Nano.

Model compression and optimization techniques are therefore crucial. A notable 2024 work by Zhu *et al.* introduced a *lightweight optimization of YOLOv3* for Jetson Nano ¹¹. They replaced YOLOv3’s backbone with an Efficient-RepVGG network optimized for GPU throughput, designed a compact bidirectional feature pyramid, and an “optimized couple head” for detection. This drastically reduced the model size – cutting ~92% of parameters and compute – and yielded a **9.95× speedup** in inference, with minimal accuracy loss ¹². The optimized YOLOv3 ran 1.6–2× faster than other detectors of similar accuracy, or conversely achieved 3–6% higher mAP at similar speed ¹³. Such architecture tailoring (e.g. using depthwise separable convolutions, smaller feature maps, etc.) has become a common strategy to **maintain accuracy while shrinking models** for edge deployment. Other research has benchmarked

EfficientDet-Lite models and SSD variants on micro UAV hardware, since those were originally designed for mobile/embedded use ¹⁴. For instance, EfficientDet-Lite0 (the smallest EfficientDet) and SSD-MobileNet were found to be very fast on devices like Pi+TPU or Orin, albeit with lower mAP than newer YOLO models ¹⁵ ⁴. **Transformer-based detectors** are generally too heavy for micro drones, but simplified versions like *RT-DETR* (Real-Time DETR) have been tested in simulation. A recent work used RT-DETR on a ground station for high-accuracy detection, while keeping a YOLOv5-Nano onboard the drone for coarse detection ¹⁶ ¹⁷. In summary, the past five years show a clear trend of developing **smaller, CNN-based detectors (YOLO nanos, MobileNet-SSD, etc.) and custom-tailored networks** that can fit the strict latency and memory constraints of onboard hardware.

Real-Time Multi-Object Tracking on Drones

Building on object detection, **multi-object tracking (MOT)** algorithms link detections across frames to track trajectories of multiple targets (e.g. people, vehicles, or animals) in view. On resource-constrained UAVs, most approaches adopt a *tracking-by-detection* paradigm: a lightweight detector processes each frame (or every Nth frame), and a tracker handles association. Simpler trackers like SORT and its variants are favored for embedded use due to their low overhead ¹⁸ ¹⁹. *SORT* (Simple Online Real-time Tracking) uses a Kalman filter and frame-to-frame data association via IoU matching, running at hundreds of FPS on a CPU ²⁰ ²¹. However, SORT can suffer frequent ID switches under occlusion. Improved versions like *DeepSORT* add a re-identification CNN to enhance robustness (re-identifying objects after brief disappearance) ²² ²³. DeepSORT's ReID model adds computational load, but can still run in real time on a Jetson-class GPU when the number of objects is modest. For example, in a ground-based test, DeepSORT with YOLOv8x was used to track vehicles and re-identify targets that the drone's onboard detector missed ¹⁷ ²⁴. Recent trackers like *ByteTrack* (which keeps low-confidence detections to improve tracking) and *OC-SORT* have also been tested on drone imagery ²⁵. In a 2025 wildlife monitoring study (WildLive), the authors benchmarked these trackers onboard and confirmed that modern MOT algorithms (OC-SORT, ByteTrack) are *computationally feasible on edge devices* while providing high accuracy ²⁶ ²⁵.

To push MOT performance for drones, researchers have proposed combining classical tracking (e.g. optical flow) with deep learning. The **WildLive** system (2025) specifically targets *onboard tracking of wildlife in high-resolution (4K) drone video* ²⁷ ²⁸. It runs on a Jetson AGX Orin and achieves near real-time performance (17+ FPS at 1080p, ~7 FPS at 4K) by using a sparse **Lucas-Kanade optical flow tracker** for most frames and only running the heavy YOLO-based detector intermittently ²⁹ ³⁰. By focusing expensive CNN inference on "regions of uncertainty" and tracking objects with cheap point trackers between detections, the system accelerates overall throughput ³⁰ ³¹. This hybrid approach yielded accurate multi-animal tracking onboard a drone, and the authors demonstrated it in field flights over a Kenyan conservancy, releasing both the **WildLive dataset** (200k+ annotated animal instances) and their source code to the community ³² ³³. Another frontier is **cross-scene MOT** for drones: Wang *et al.* (2023) tackled the challenge of a UAV tracking targets through highly dynamic scenarios (varying altitude, motion, etc.) ³⁴ ³⁵. They proposed a meta-learning enhanced tracker that can adapt to new scenes on the fly, and a re-identification method fusing camera and drone state (IMU/GPS) data via Dempster-Shafer theory ³⁶ ³⁷. Their framework included a "*lightweight detection module*" with attention mechanisms to handle small aerial targets ³⁸, and a trajectory predictor using an LSTM improved by Model-Agnostic Meta-Learning ³⁹ ⁴⁰. A new **MIDDDTD dataset** (Multi-Info Drone Detection and Tracking) was introduced with diverse scenes and drone metadata to evaluate such cross-scene trackers ⁴¹ ⁴². This indicates a growing interest in **domain-specific MOT solutions for drones**, although many of these advanced techniques are still tested offline or on high-end hardware. On actual micro drones, simpler real-time tracking-by-detection (possibly augmented with occasional ReID or flow tracking) remains the pragmatic choice.

Embedded Hardware and Acceleration Techniques

Embedded AI hardware has rapidly improved in recent years, enabling more complex vision models to run onboard small drones. NVIDIA's Jetson series (Nano, TX2, Xavier NX, Orin Nano, etc.) provides GPU acceleration in a compact form factor, while alternatives like the Raspberry Pi rely on CPU or external accelerators (Google Coral Edge TPU, Intel Movidius NCS2) for neural network inference. A 2024 review noted that the Jetson Xavier NX outperforms Jetson Nano and Raspberry Pi 4 + NCS2 setups in real-time object detection workloads ⁴³. In fact, the Jetson Xavier NX can achieve ~30 FPS with a YOLOv7-tiny model, versus single-digit FPS on the Nano or Pi ³. The **Jetson Orin Nano (2023)** has further closed the gap, delivering up to 20 TOPS of AI performance. As mentioned, tests show it can run modern detectors at tens of FPS with modest power draw ⁴ ⁴⁴. Among CPU-based platforms, the **Raspberry Pi 5 (2023)** improved on its predecessors but still only reaches a few FPS on small YOLO models without acceleration ⁴⁵ ⁴⁶. Thus, **neural accelerators** are essential on Pi-class boards: the Coral Edge TPU, for example, can run quantized MobileNet-SSD detectors at over 100 FPS in ideal conditions ⁴⁷ (though real-world performance is lower due to I/O overhead). One study demonstrated YOLOv5s running on a Coral USB accelerator attached to a Pi, but noted the USB2 vs USB3 connection significantly throttles throughput (USB2 incurred a 3× slowdown) ⁴⁸ ⁴⁹. Overall, leveraging hardware like the **Edge TPU, NCS2, or GPU Tensor Cores** is critical for achieving real-time inference on small drones.

In tandem, researchers apply a variety of **model optimization techniques** to maximize inference speed on constrained hardware. **Quantization** is widely used: converting 32-bit floating point networks to 8-bit integer precision can dramatically speed up inference on supported hardware (TPUs, many DSPs) and reduce memory usage. The trade-off is slight accuracy degradation if not carefully done ⁵⁰ ⁵¹. Many UAV vision models now use INT8 or FP16 weights. For instance, TakuNet (2024) – a tiny CNN for drone image classification – was trained with FP16 and deploys via NVIDIA TensorRT to exploit low-bit operations, achieving an impressive 650 FPS on a Jetson Orin Nano (15 W) ⁵². **Pruning and weight compression** are also explored: redundant channels/filters are pruned to lighten models ⁵⁰. This must be balanced to avoid too much accuracy loss. The optimized YOLOv3 by Zhu *et al.* effectively pruned YOLOv3 by replacing layers and heads, yielding a model an order of magnitude faster on Jetson Nano with minimal accuracy drop ⁵³. Another tactic is using **neural architecture search (NAS)** or manual architecture redesign to create efficient models. We've seen efficient backbones (MobileNetV3, ShuffleNet, etc.) and novel blocks (e.g. CSP, depthwise conv) adopted in nearly all recent detectors ⁵⁴ ⁵⁵. Many of these ideas were crucial for producing lightweight drone-ready networks. Additionally, **inference optimizers and libraries** play a big role on embedded devices. NVIDIA's TensorRT is frequently used to compile models to optimized GPU kernels, often doubling FPS compared to naive execution ⁵⁶ ⁵⁷. For CPUs, frameworks like TensorFlow Lite and OpenVINO are used to leverage NEON instructions and multi-threading. In a late-2024 benchmark, Alqahtani *et al.* deployed various detectors (YOLOv8-n/s/m, EfficientDet-Lite0-2, SSD) across devices using TensorRT on Orin Nano and TFLite on Raspberry Pi/TPU, carefully measuring energy and latency ¹⁴ ⁴. They found the Orin Nano in TensorRT mode offered the best overall speed/accuracy/energy balance – e.g. only 16 ms inference for YOLOv8-n and ~20 ms for EfficientDet-Lite on Orin ¹⁵ ⁴ – while the Pi 4 with Edge TPU achieved faster inference per watt on some models but at a cost of lower accuracy due to aggressive quantization ⁴⁴ ⁵⁸. This underscores that **system-level optimization (choosing the right model, precision, and runtime for the hardware)** is as important as the model design itself.

Onboard Perception-Control Integration

A critical aspect of these advances is how detection and tracking results are integrated into the **drone's control and navigation system**. Running perception onboard only becomes useful if the drone can autonomously react to the identified objects in real time. Many researchers leverage the Robot

Operating System (ROS) or similar middleware to connect vision outputs to flight controllers. For example, object detection packages running on a Jetson can publish target coordinates to a PX4 flight controller to implement “follow that object” behaviors (community projects have demonstrated a Jetson Nano controlling a drone to track a selected object via live video ⁵⁹). In the academic works reviewed, several explicitly close the loop between perception and control. The *Dual-Stage UAV Processing Architecture* by Ntousis *et al.* (2025) uses an onboard RPi+NCS2 to detect vehicles (via YOLOv5n) and immediately adjusts the drone’s heading towards the detected target, while a ground server handles more advanced trajectory planning ¹⁶ ⁶⁰. Their system maintained a high-speed link to offload heavy computations, but the initial detection and guidance commands were generated onboard in real time. Similarly, the WildLive system aims for **animal-reactive flight**: the eventual goal is for the UAV to autonomously adjust its path to keep animals in view or avoid startling them, using the onboard detection/tracking feed ²⁷ ⁶¹. They note this capability is essential for beyond-visual-line-of-sight wildlife monitoring, where relying on a remote pilot or data link is impractical ²⁷ ⁶². Initial tests of WildLive were run with manual supervision, but it demonstrated that the drone could *onboardly* identify multiple animals and their trajectories in a large 4K scene ²⁹ ²⁵, which is a step toward closed-loop animal-following drones.

In many cases, **state estimation and sensor fusion** are integrated with vision to improve robustness. Drones have onboard IMUs, GPS, altitude sensors, etc. Some recent trackers (e.g. Wang *et al.* 2023) explicitly incorporate drone state into the tracking algorithm to help predict target motion relative to the moving camera ⁶³ ³⁷. This hints at a broader integration where the vision module might send *observations* to the flight control stack, which fuses them with inertial navigation data. For instance, a detected object’s bearing and range (if depth is available or via size cues) could be treated as a measurement in the drone’s EKF, enabling target localization and pursuit. There have also been demonstrations of visual odometry and SLAM running alongside object detection on micro drones, so that obstacles/objects are mapped and avoided in real time. The tight coupling of **perception and control** is facilitated by middleware and frameworks. DJI’s SDK and autonomous drone platforms (like the Parrot ANAFI Ai) have started to allow running custom neural networks onboard, feeding the results into their guidance systems. On the research side, open-source frameworks like **AirSim** and **Flightmare** (a simulator from 2021) have been used to simulate how a drone’s control responds to onboard vision outputs, before real flight trials. Overall, the last five years show a move from using onboard vision purely for data collection to using it for **closed-loop autonomy**, where the drone can adjust its flight path on the fly based on what it sees. Achieving low-latency processing (often requiring the optimizations discussed earlier) is key to such integration, so that control decisions (e.g. to avoid an object or track a person) are made quickly enough for stable flight.

Benchmarks and Comparisons

To quantify progress, researchers have published **benchmark studies** comparing models and hardware for drone vision. One comprehensive benchmark is **YOLOBench (2023)**, which evaluated 550+ YOLO-family models across 4 different embedded hardware setups (x86 CPU, ARM CPU, Nvidia Jetson GPU, Edge TPU) ⁶⁴ ⁶⁵. By measuring accuracy (mAP) vs latency for each model, they constructed Pareto-optimal frontiers for different platforms. An interesting finding was that **older YOLO models (e.g. YOLOv3, YOLOv4)**, when retrained with modern techniques, can be as efficient as newer models on certain edge hardware ⁶⁶. For example, on a Raspberry Pi 4 CPU, a pruned-and-quantized YOLOv4 might achieve a better accuracy-speed trade-off than an out-of-the-box YOLOv8, due to the latter’s greater complexity ⁶⁵ ⁶⁷. Another benchmark (Santos *et al.* 2024) specifically measured *YOLOv7-tiny* on three devices, as noted earlier: **RPi 4 (0.9 FPS), Jetson Nano (~6 FPS), Xavier NX (30 FPS)** ³. They also observed that Jetson Nano’s two power modes make a difference (max-N mode yielded ~7.4 FPS vs ~5.2 FPS in low-power mode) ³. A related 2024 study tested *YOLOv7-tiny with TensorRT* on the Jetson Xavier NX and achieved the full 30 FPS at just 10 W power, which demonstrates how much optimization

can improve edge performance ⁶⁸ ⁶⁹. For Raspberry Pi-based setups, benchmarks indicate that without offloading, high-performance real-time detection is infeasible – hence the reliance on USB accelerators. A comparison of the **Intel Movidius NCS2 vs Google Coral TPU** on a Raspberry Pi 4 showed that the Coral can reach ~15–20 FPS on MobileNet-SSD at 300×300 resolution, about 3× faster than the NCS2 on the same task ⁴⁷ ⁷⁰. However, the Coral’s requirement of 8-bit quantization meant the accuracy on COCO dropped a few points compared to FP16 models on Jetson ⁴⁴ ⁷¹.

Table 1 summarizes some performance metrics from recent works, illustrating the range of trade-offs:

Model & Hardware	Accuracy (mAP)	Speed (FPS)	Notes
YOLOv7-tiny on Raspberry Pi 4 CPU	53% (COCO val)	0.9 FPS ³	No accelerator; very slow.
YOLOv7-tiny on Jetson Nano	53% (COCO val)	~6 FPS ³	~7 FPS in 10 W mode; limited by 128 CUDA cores.
YOLOv7-tiny on Jetson Xavier NX	53% (COCO val)	30 FPS ³	Real-time achieved with TensorRT optimizations.
YOLOv8-Nano on Jetson Orin Nano	~40% (COCO val)	60 FPS ⁴	FP16 TensorRT engine; ~16 ms inference.
EfficientDet-Lite0 on Jetson Orin Nano	~32% (COCO val)	50 FPS ⁴	Small TPU-optimized model, GPU TensorRT used.
YOLOv5n on Raspberry Pi + NCS2	~45% (COCO val)	5–8 FPS (est.)	Used in dual-stage system ¹⁶ (vehicles).
YOLOv5s on Raspberry Pi + Coral TPU	~47% (COCO val)	16 FPS ⁴⁸ ⁴⁹	USB3 interface; INT8 quantized model.
CenterNet (ResNet18) on Jetson Nano	33% (VisDrone) ⁷² ⁷³	4.5 FPS ⁹	Small-object focused detector (Drones 2023) ² .
WildLive (YOLO+flow) on Jetson Orin	85% (wildlife dataset)	17 FPS (1080p) ²⁸ ³⁰	Hybrid tracking mode; 7 FPS at 4K resolution.

Table 1: Performance examples of object detection/tracking models on embedded platforms. Accuracy figures are approximate for context (different datasets), and FPS measured under specific conditions per cited sources.

These comparisons show that **real-time (≥ 30 FPS) object detection on micro UAVs is now achievable** on higher-end embedded boards (Xavier NX, Orin), while lower-end setups (Nano, Pi) typically reach 5–15 FPS without offloading. Techniques like quantization and TensorRT provide significant boosts, often 2–3× speedups ¹¹ ¹³. It is also evident that there is no one-size-fits-all solution: the best model depends on the hardware and the task. For instance, a tiny model might suffice for obstacle avoidance (where ~60% mAP on simple objects is fine) but a surveillance drone might require a larger model for reliable person detection, accepting lower FPS or using a stronger onboard computer.

Open Source Implementations and Datasets

The rapid progress in this field has been supported by the release of **open-source code and datasets** tailored to drones. The **VisDrone dataset** (2018–2021) and **UAVDT** provided thousands of annotated

images and videos of urban scenes from drones, spurring research into aerial pedestrian and vehicle detection. By 2020, VisDrone’s challenge had a dedicated MOT track, and many of the lightweight models reviewed (e.g. YOLO variants, CenterNet adaptations) report results on it ⁷⁴ ⁷⁵. More recently, specialized datasets have emerged: the **AU-AIR dataset** (2020) focuses on UAV imagery with objects annotated for autonomous navigation scenarios; the **Drone-Crowd** dataset (2021) provides dense annotations for crowd counting and tiny person detection from drones. In 2023–2024, as mentioned, **WildLive** released a large wildlife tracking dataset ³², and **MIDDDTD** (2023) was made available with multi-modal data for drone tracking ⁴¹ ⁴². These datasets enable researchers to train and evaluate models under realistic conditions (e.g., animals in savannah terrain for WildLive, or mixed urban-rural scenes with drone telemetry in MIDDDTD).

On the implementation side, many researchers have shared their frameworks. The WildLive project provides all source code and trained model weights for its YOLO-based detection and LK tracking framework, accessible on GitHub ³³. This allows the community to reproduce and build upon their near-real-time tracking results. The YOLOBench benchmark similarly released its code and the 550+ model zoo for others to analyze or use in NAS research ⁶⁷. Popular open-source repositories like **Ultralytics YOLOv5/YOLOv8** (which are frequently cited in papers) have made it easy to obtain and modify state-of-the-art detectors for custom drone tasks. There are also ROS packages integrating these detectors, and drone control software (e.g., **DJI UX SDK**, **MAVSDK**) with examples of using onboard vision. A notable open project is the **search-and-rescue drone with Coral TPU** on Hackster.io ⁷⁶, which provides a step-by-step implementation of a drone using a Pi + Coral to detect and localize missing persons. Academic efforts like the dual-stage architecture ¹⁶ ⁷⁷ and cross-scene tracker have pseudocode or algorithm details in their papers that can guide implementation. As hardware becomes more accessible, we also see **community benchmarks** (on forums and GitHub) comparing, say, TensorFlow Lite vs. TensorRT on the same Jetson, or testing new models like YOLOv8 on the Jetson Nano – these informal reports often complement academic studies and help practitioners choose the right setup.

In summary, the ecosystem of **datasets and open implementations** has matured from general-purpose to drone-specific. This accelerates development of better models: for instance, having aerial datasets allowed researchers to fine-tune detectors for the unique appearance of objects from a bird’s-eye view (small, occluded, different angles) ⁷⁸ ⁷⁹. Open benchmarks and code ensure that claims of “real-time” truly translate on the target hardware, and enable continued improvements. The community’s commitment to open science in this domain is evident from the number of releases in the last few years.

Conclusion

Efficient real-time vision on micro aerial robots has advanced markedly since 2020. Researchers have devised **ultra-lightweight object detectors**, from tiny YOLO variants to custom networks, that push the limits of what can run on a few watts. Multi-object tracking approaches have been tailored for UAV constraints, often by smartly reducing workload (e.g. combining fast optical flow with intermittent CNN detection). These perception modules are increasingly being integrated with **onboard flight control**, moving drones closer to true autonomy in complex missions. While a sub-250g quadrotor may still struggle to carry hardware for high-resolution video analytics at high FPS, the trend of more powerful yet efficient edge AI devices (like the Jetson Orin Nano and upcoming Qualcomm drone AI SOC) is rapidly closing that gap. The latest works demonstrate that even tasks like 4K animal tracking or on-drone pedestrian detection are within reach of onboard processing ³⁰ ⁸⁰. Continued innovation in model compression, hardware acceleration (including dedicated AI chips for drones), and cooperative architectures (onboard vs offboard processing) will further improve performance.

Crucially, the **real-world deployments** reviewed – from wildlife conservation drones to search-and-rescue prototypes – show that these research advances are not just academic. There are now open-source reference designs and datasets that pave the way for engineers to deploy object detection and MOT on their own UAV platforms. In the coming years, we can expect micro drones to reliably detect and track multiple objects in real time, enabling applications like autonomous inspection, tracking moving subjects for filming, and smarter collision avoidance. The progress from 2020 to 2025 in real-time onboard vision is laying the groundwork for the next generation of fully autonomous micro aerial robots.

Sources:

- Lago, Patel & Singh (2024) – *Low-cost real-time aerial object detection and tracking pipeline* 39 38
 - Hossain & Lee (2019) – *Onboard multi-object detection/tracking on UAV with GPU* 81
 - Zhu et al. (2024) – *Lightweight optimization of YOLOv3 on Jetson Nano* 11 13
 - Bhowmik et al. (2023) – *On-board small-scale object detection on Jetson & NCS2* 2 9
 - Nguyen et al. (2025) – *WildLive: real-time animal detection/tracking on UAV (Jetson Orin)* 30 25
 - Wang et al. (2023) – *Cross-scene MOT for drones with meta-learning (MIDDDTD dataset)* 41 40
 - Ntousis et al. (2025) – *Dual-stage UAV detection (RPi+NCS2 onboard & YOLOv8x+DeepSORT offboard)* 16 77
 - Santos et al. (2024) – *YOLOv7-tiny performance on edge devices (Raspberry Pi, Jetson)* 3
 - Alqahtani et al. (2024) – *Benchmark of YOLOv8, EfficientDet, SSD on Pi 3/4/5 + TPU, and Jetson Orin* 15 4
 - Lazarevich et al. (2023) – *YOLOBench: 550+ models benchmark on embedded systems* 64 66 .
-

1 3 11 12 13 43 53 68 69 Performance analysis of real-time object detection on Jetson device | Request PDF

https://www.researchgate.net/publication/363515934_Performance_analysis_of_real-time_object_detection_on_Jetson_device

2 5 6 7 8 9 10 72 73 74 75 78 On-Board Small-Scale Object Detection for Unmanned Aerial Vehicles (UAVs)

<https://e-archivo.uc3m.es/rest/api/core/bitstreams/27b16e8c-e6bb-46ae-b0b3-de1efa5bae0a/content>

4 14 15 44 45 46 58 71 Benchmarking Deep Learning Models for Object Detection on Edge Computing Devices

<https://arxiv.org/html/2409.16808v1>

16 17 24 60 77 A Dual-Stage Processing Architecture for Unmanned Aerial Vehicle Object Detection and Tracking Using Lightweight Onboard and Ground Server Computations

<https://www.mdpi.com/2227-7080/13/1/35>

18 19 25 26 27 28 29 30 31 32 33 61 62 WildLive: Near Real-time Visual Wildlife Tracking onboard UAVs

<https://arxiv.org/html/2504.10165v1>

20 21 22 23 34 35 36 37 38 39 40 41 42 63 79 Cross-Scene Multi-Object Tracking for Drones: Leveraging Meta-Learning and Onboard Parameters with the New MIDDTD

<https://www.mdpi.com/2504-446X/9/5/341>

47 Models - Object Detection - Coral

<https://coral.ai/models/object-detection/>

48 49 70 riverpublishers.com

https://www.riverpublishers.com/downloadchapter.php?file=RP_9788770227902C11.pdf

50 51 52 54 55 56 57 TakuNet: an Energy-Efficient CNN for Real-Time Inference on Embedded UAV systems in Emergency Response Scenarios

<https://arxiv.org/html/2501.05880v2>

59 Object tracking system - Jetson Nano - NVIDIA Developer Forums

<https://forums.developer.nvidia.com/t/object-tracking-system/178600>

64 65 66 67 [2307.13901] YOLOBench: Benchmarking Efficient Object Detectors on Embedded Systems

<https://arxiv.org/abs/2307.13901>

76 Search And Rescue Drone With Google Coral - Hackster.io

<https://www.hackster.io/bandofpv/search-and-rescue-drone-with-google-coral-a485c7>

80 Towards Real-Time On-Drone Pedestrian Tracking in 4K Inputs

<https://www.mdpi.com/2504-446X/7/10/623>

81 Deep Learning-Based Real-Time Multiple-Object Detection ... - MDPI

<https://www.mdpi.com/1424-8220/19/15/3371>